



# Dev-MongoDB Best Practices About Index

Kuririn Wade

索引是



索引是



# 类别

- Index
- Unique Index
- Sub-document index
- Compound Index
- ...

# 类别

- Index
- Unique Index
- Sub-document index
- Compound Index
- ...

Index Options



|   | Index Type                | Description   |
|---|---------------------------|---|
| 1 | Primary Index             | Index on the document key on the whole bucket   |
| 2 | Named Primary Index       | Give name for the primary index. Allows multiple primary indexes in the cluster   |
| 3 | Secondary Index           | Index on the key-value or document-key  |
| 4 | Secondary Composite Index | Index on more than one key-value  |
| 5 | Functional Index          | Index on function or expression on key-values   |
| 6 | Array Index               | Index individual elements of the arrays   |
| 7 | Partial Index             | Index subset of items in the bucket   |
| 8 | Covering Index            | Query able to answer using the the data from the index and skips retrieving the item.   |
| 9 | Duplicate Index           | This is not type of index. Feature of indexing that allows load balancing. Thus providing scale-out, multi-dimensional scaling, performance, and high availability. |

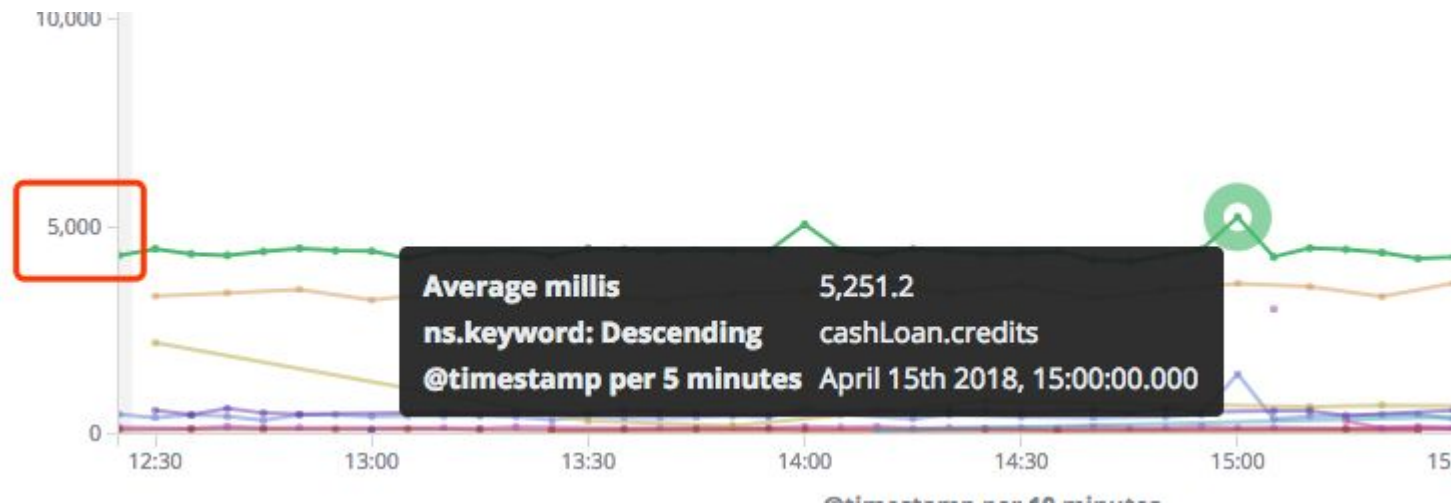
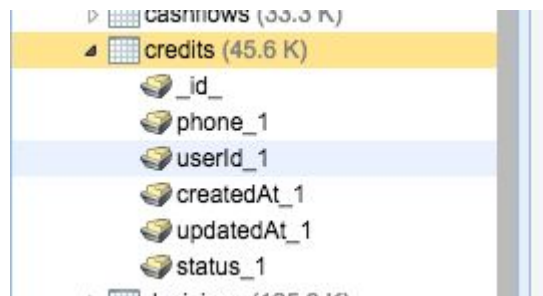
# 类别: Compound Index

- Difference on Compound index order {userId:1,date:1} vs {date:1,userId:1}
- Difference on Sort order {userId:-1} vs {userId:1}
- Should I create two separate indexes or compound index or both? -- [It depends](#)

# 典型的慢查询

```
"cursorExhausted": "True",
"nreturned": 0,
"numYield": 6000,
"ns": "new_bear.user",
"execStats": {
  "invalidates": 0,
  "works": 768110,
  "executionTimeMillisEstimate": 800,
  "restoreState": 6000,
  "stage": "LIMIT",
  "needTime": 768109,
  "nReturned": 0,
  "saveState": 6000,
  "inputStage": {
    "works": 768110,
    "executionTimeMillisEstimate": 800,
    "restoreState": 6000,
    "stage": "COLLSCAN",
    "needTime": 768109,
    "needYield": 0,
    "nReturned": 0,
    "filter": {
      "id": {
        "$eq": "599453884294fd9d6c88a13c"
      }
    },
    "saveState": 6000,
    "advanced": 0,
    "isEOF": 1,
    "direction": "Forward",
    "invalidates": 0,
    "docsExamined": 768108
  },
  "isEOF": 1,
  "needYield": 0,
  "limitAmount": 1,
  "advanced": 0
}
```

# 典型的慢查询





# 典型的慢查询

```
docsExamined:13668,
nreturned:0,
execStats:[] {
  limitAmount:5,
  advanced:0,
  invalidates:0,
  executionTimeMillisEstimate:4679,
  needTime:13668,
  isEOF:1,
  saveState:236,
  needWrite:0,
  stage: LIMIT,
  restoreState:236,
  nReturned:0,
  inputStage:[] {
    nReturned:0,
    advanced:0,
    saveState:236,
    transformBy:[] {
      userId:userId
    }
  },
  executionTimeMillisEstimate:4679,
  needTime:13668,
  isEOF:1,
  invalidates:0,
  needYield:0,
  stage: PROJECTION,
  restoreState:236,
  inputStage:[] {
    nReturned:0,
    filter:[] {
      $and:[] {
        [] {
          report.mnoCommonlyConnectMobiles:[] {
            $exists:True
          }
        },
        [] {
          $nor:[] {
            [] {
              report.mnoCommonlyConnectMobiles.mobile:[] {
                $exists:True
              }
            }
          }
        }
      }
    }
  }
}
```

```
inputStage:[] {
  direction:forward,
  saveState:236,
  keysExamined:13668,
  dupsTested:0,
  multiKeyPaths:[] {
    status:[] {
    }
  },
  isEOF:1,
  executionTimeMillisEstimate:20,
  needTime:0,
  isSparse:False,
  seenInvalidated:0,
  seeks:1,
  isUnique:False,
  advanced:13668,
  invalidates:0,
  isPartial:False,
  indexBounds:[] {
    status:[] {
      [] {
        normal,
        normal
      }
    }
  },
  restoreState:236,
  dupsDropped:0,
  indexName:status_1,
  isMultiKey:False,
  needYield:0,
  stage: IXSCAN,
  indexVersion:0,
  keyPattern:[] {
    status:1
  },
  nReturned:13668,
  works:13669
},
works:13669
},
works:13669
},
works:13669
```

# 典型的慢查询

Usually building index on status can actually HARM on performance

# 理想的查询

- IXSCAN
- Limit
- Projection
- Docs Size less is better
- Docs Examined less is better
- $nReturned = totalDocsExamined = totalKeysExamined$
- SORT in index
- Limit Array elements

# 理想的查询

对于普通查询, 我们最希望看到的组合有这些:


- Fetch+IDHACK
- Fetch+ixscan
- Limit+(Fetch+ixscan)
- PROJECTION+ixscan
- SHARDING\_FILTER+ixscan
- Sort + Index

# 理想的查询

- SSD + Large RAM
- Consider sharding when data size > RAM

# 理想的查询

**Ensure indexes fit in RAM**



# 理想的查询

不希望看到包含如下的stage:

- COLLSCAN(全表扫), SORT(使用sort但是无index), 不合理的SKIP, SUBPLA(未用到index的\$or)

对于count查询, 希望看到的有:

- COUNT\_SCAN

不希望看到的有:

- COUNTSCAN

## Explain() 指标

### Explain

```
db.collection.find(query).explain();

{
  "cursor" : "BasicCursor",
  "indexBounds" : [ ],
  "nscanned" : 57594,
  "nscannedObjects" : 57594,
  "n" : 3 ,
  "millis" : 108
}
```



# Explain() 指标

```
"executionStats" : {  
  "executionSuccess" : true,  
  "nReturned" : 2,  
  "executionTimeMillis" : 0,  
  "totalKeysExamined" : 4,  
  "totalDocsExamined" : 2,  
  "executionStages" : {  
    "stage" : "FETCH",  
    "nReturned" : 2,  
    ...  
    "inputStage" : {  
      "stage" : "IXSCAN",  
      "nReturned" : 2,  
      ...  
      "keyPattern" : {  
        "a" : 1,  
        "c" : 1,  
        "b" : 1  
      },  
      "indexName" : "a_1_c_1_b_1",  
      "isMultiKey" : false,  
      "direction" : "backward",  
      "indexBounds" : {  
        "a" : [  
          "[1.0, 1.0]"  
        ],  
        "c" : [  
          "[MaxKey, MinKey]"  
        ],  
        "b" : [  
          "(3.0, -inf.0]"  
        ]  
      },  
      "keysExamined" : 4,  
      "dupsTested" : 0,  
      "dupsDropped" : 0,  
      "seenInvalidated" : 0,  
      "matchTested" : 0  
    }  
  }  
}
```

# Explain() 指标

[more explain](#)

# 指标

[More in cloud manager index](#)

# So, 越多索引越好？

- According to Use frequency
- Index Consume space
- Update index after new data insert
- Consume time during restore data

# 坑

- Foreground Build Index (default) will block operation
- Upsert won't **duplicate** ONLY when in unique index
- Avoid array length too large (performance & 16M Documents size & Agg Limit)

# What If

- Q:我觉得我的DB优化到头了, 降不下来query 时间?

# What If

Answer:

- 我觉得
- Cache
- Hardware(Mem,Sharding,...,Enterprise Version)
- 历史数据Snapshot

# 经验

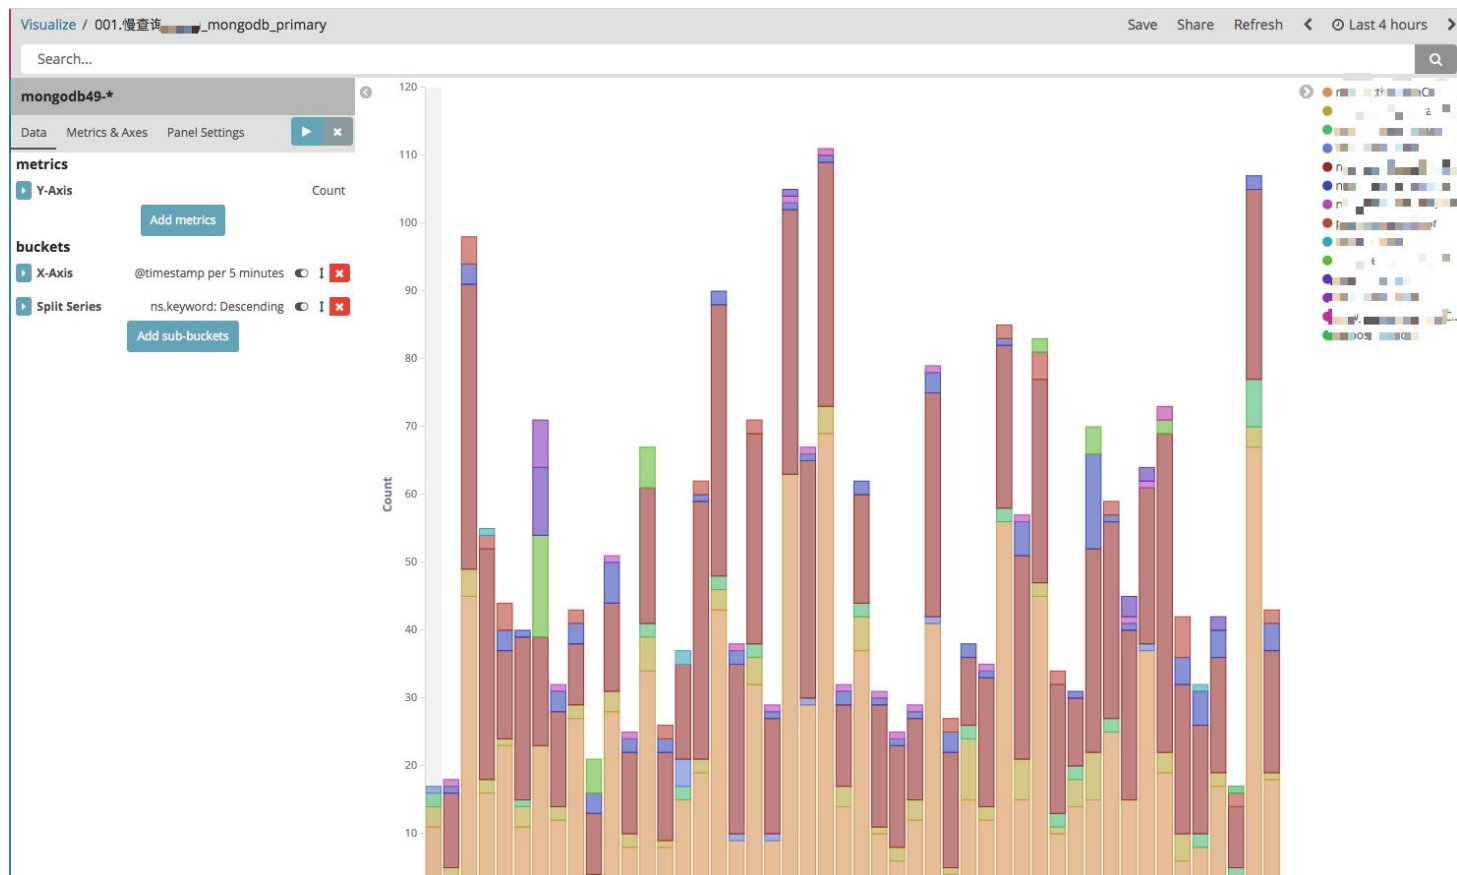
- `count()` > 10K must create index
- Query should < 10ms
- Index should < 100 ms




# 无数据不优化

```
db.setProfilingLevel(1,100)
```

# 无数据不优化



# 无数据不优化



# More Tuning(not only MongoDB)

## N1QL Access Methods and Performance



Fastest to slowest, 1 to 5

|   | Method                    | Description   |
|---|---------------------------|---|
| 1 | <b>USE KEYS</b>           | Document fetch, no index scan   |
| 2 | <b>COVERED Index Scan</b> | Query is (or part of the query during JOIN) is processed with index scan only |
| 3 | <b>Index Scan</b>         | Partial index scan, then fetches  |
| 4 | <b>JOIN</b>               | Fetch of left-hand-side, then fetches of right-hand-side                      |
| 5 | <b>Primary Scan</b>       | Full bucket scan, then fetches  |



# More Tuning(not only MongoDB)

## Advice on Query Performance



- EXPLAIN to analyze query plan
- Index selection, spans for push down of as many predicates as possible. More the merrier
- Pushdown of LIMIT, OFFSET
- Index order for ORDER BY
- Covering index
- Simple COUNT queries can take advantage of index count
- Exploit index for MIN queries
- For ANY, ANY AND EVERY, WITHIN predicates use ARRAY index.
- For UNNEST, use ARRAY index. Array key has to be the leading key (Only for UNNEST)
- USE IN instead of WITHIN
- Use pretty=false (4.5.1), max\_parallelism when queries return large resultset
- Improve fetch performance by increasing pipeline-cap, pipeline-batch
  - Exploit array fetch by query rewrite
- Execute query and explore each phase of monitoring stats of query.
- Monitor CPU and memory usage and adjust number of Query Service Nodes.



# More Tuning

<https://www.slideshare.net/mongodb/fast-querying-indexing-for-performance-4> #54

<https://www.slideshare.net/journalofinformix/tuning-for-performance-indexes-queries>

<https://www.slideshare.net/mongodb/indexing-with-mongodb>

<https://www.slideshare.net/mongodb/indexing-and-query-optimizer-richard-kreuter>

<https://www.slideshare.net/mdirolf/indexing>